

# Pearl

Tor relay implementation in Golang  
31 October 2017

Michael McLoughlin  
Software Engineer, Uber

# Background


Motivations:

- Learn about the Tor protocol
- Learn by doing!
- Goaded by a glib writeup of a [previous Go implementation](#)
- "After a while I realized the language was not suited for the project"


Thus pearl was born out of *hubris* on my part.

- Bradfield Sabbatical Program was a great way to revive my project
- Turns out it's non-trivial

# Tor network enables anonymous internet access




## Clients send traffic via a circuit




3-hop circuits consist of: **guard** (red), **middle** (green) and **exit**.

## Circuit initiated with a CREATE command




Shared secrets established for symmetric encryption.

## Circuit extended with RELAY\_EXTEND




Relay to *green* node without green knowing the gopher's location.

Repeated to complete 3-hop circuit




Now streams can be established on the circuit.

Data can now be relayed







Each node removes its layer of encryption




It's like peeling an onion...



Exit node sees the actual packet




## The crux of a relay implementation




Relay maintains many server and client connections.

Traffic on a circuit is relayed between them



## Demux/mux circuit traffic



Relay must demux/mux circuit traffic. Circuits must not interfere with each other.

# Why Go?

Strikes a sweet spot

- Fast statically typed compiled language
- Feels like writing in a dynamic language
- Strong tooling, standard library and ecosystem

Concurrency features:

- **Goroutines:** "green threads" mapped to OS threads, running ~100k is reasonable
- **Channels:** send and receive operations between goroutines

Gained a reputation for:

- Systems programming
- High performance servers

## Why not?

Tom van der Woerd had [harsh criticisms](#):

- Limited cipher suites and poor performance of crypto/tls
- Hence forced to use C bindings to openssl
- Blamed cgo for buildup of OS threads and excessive locking
- Memory usage per connection: 16KB buffer plus 4KB goroutine stack
- Buffered channel: static array implementation causes high memory usage

Maybe we can avoid these problems now:

- crypto/tls has massively improved
- Memory problems result from *design decisions*



# Challenges

Grunt work:

- Forking crypto/tls standard library package
- Parsing protocol data formats
- Navigating a *work-in-progress* spec

Cryptographic details:

- Auth flow
- NTor handshake
- Old TAP handshake
- Some hand-rolled algorithms

# Concurrency is hard

State managed by an associated goroutine

- Connection
- TransverseCircuit

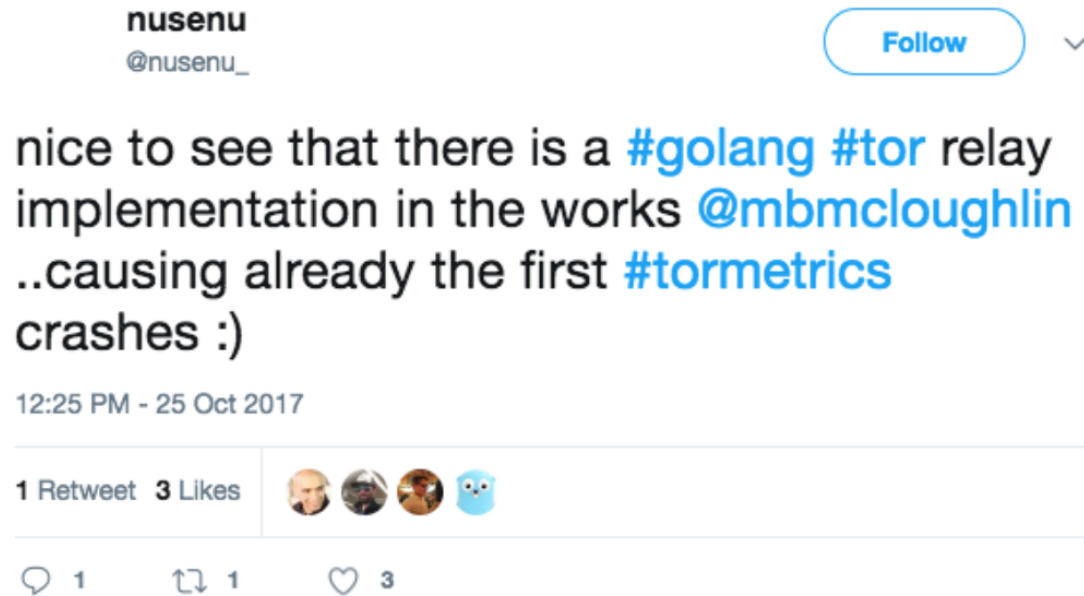
Connection offloads to circuit via channel:

```
cell, _ := c.ReceiveCell()
switch cell.Command() {
case CommandCreate, CommandCreate2:
    CreateHandler(c, cell)
case CommandCreated, CommandCreated2, CommandRelay, CommandRelayEarly, CommandDestroy:
    s, ok := c.circuits.Sender(cell.CircID())
    if !ok {
        bad()
    }
    s.SendCell(cell)
case CommandPadding, CommandVpadding:
    logger.Debug("skipping padding cell")
default:
    logger.Error("no handler registered")
}
```

## Current Status



# It's gone viral

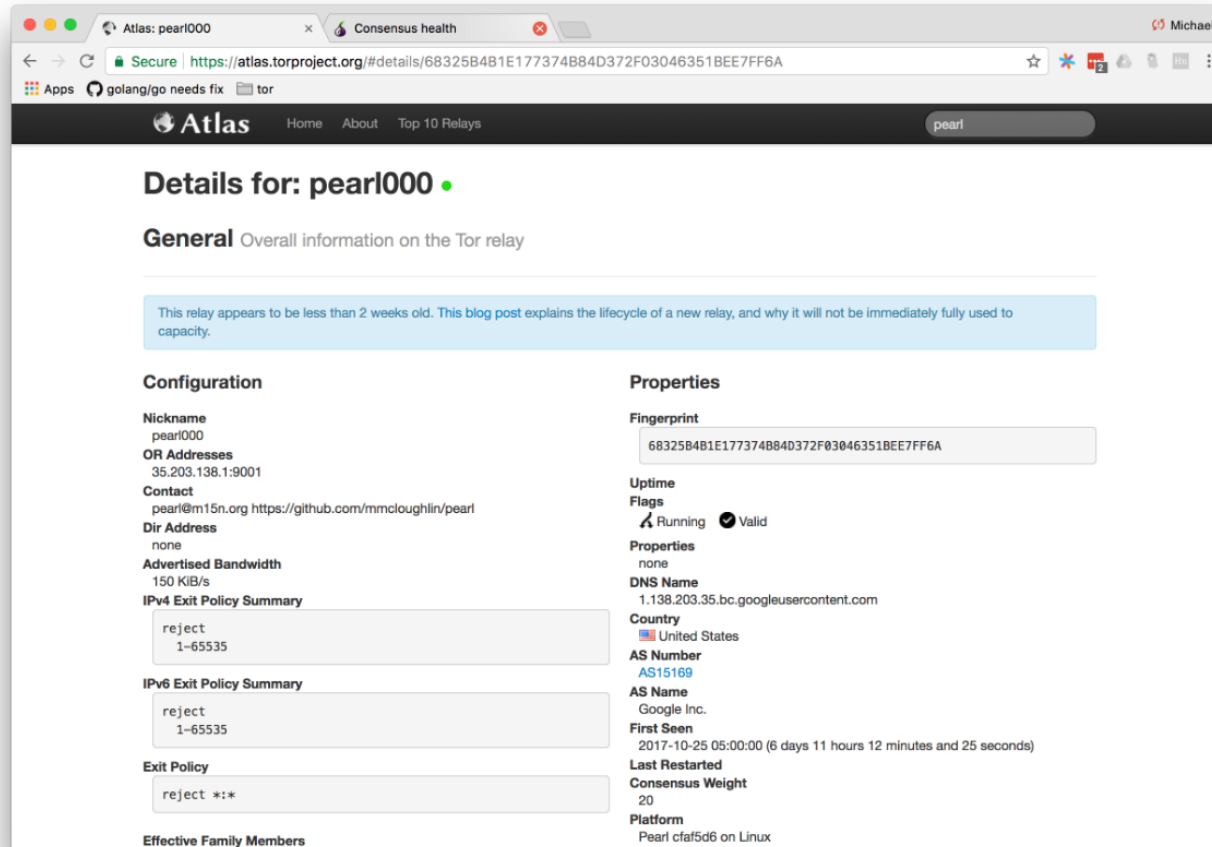


(Not quite. I did crash the tor metrics system though.)

# Accepted into consensus

```
$ wget http://86.59.21.38/tor/server/all
$ cat all
router pearl000 35.203.138.1 9001 0 0
signing-key
-----BEGIN RSA PUBLIC KEY-----
MIGJAoGBAKzTaN4tZGv1kiQWBzeuOk+ovr2LtIURlaVC38j6j/fQuYfuAZX/XvV1
fQr9EVh+T617dh+frr2D0QDuzLUvP3hpgVozW94w+Ib85pUCne03f4rj3QYu5Qtg
GvzShslZI6vggy0g2jA0Ga4jxT/UYAcKE5dQo8CBKA6Qb0P5Joc1AgMBAAE=
-----END RSA PUBLIC KEY-----
fingerprint 6832 5B4B 1E17 7374 B84D 372F 0304 6351 BEE7 FF6A
...
platform Pearl cfaf5d6 on Linux
contact pearl@m15n.org https://github.com/mmcloughlin/pearl
...
```

# Pearl node running in the Tor network



The screenshot shows a web browser window with two tabs: 'Atlas: pearl000' and 'Consensus health'. The address bar shows a secure connection to <https://atlas.torproject.org/#details/68325B4B1E177374B84D372F03046351BEE7FF6A>. The page title is 'Details for: pearl000'. The 'General' section provides overall information on the Tor relay, including a note that the relay is less than 2 weeks old. The 'Configuration' section lists the nickname 'pearl000', OR addresses '35.203.138.1:9001', contact 'pearl@m15n.org', and advertised bandwidth of '150 KiB/s'. The 'Exit Policy Summary' for both IPv4 and IPv6 is 'reject 1-65535'. The 'Exit Policy' is 'reject \*:\*'. The 'Effective Family Members' section is empty. The 'Properties' section shows the fingerprint '68325B4B1E177374B84D372F03046351BEE7FF6A', uptime, flags 'Running' and 'Valid', and DNS name '1.138.203.35.bc.googleusercontent.com'. The 'Country' is 'United States', 'AS Number' is 'AS15169', 'AS Name' is 'Google Inc.', 'First Seen' is '2017-10-25 05:00:00 (6 days 11 hours 12 minutes and 25 seconds)', 'Last Restarted' is '2017-10-25 05:00:00 (6 days 11 hours 12 minutes and 25 seconds)', 'Consensus Weight' is '20', and 'Platform' is 'Pearl cfa5d6 on Linux'.

Atlas Home About Top 10 Relays pearl

## Details for: pearl000

**General** Overall information on the Tor relay

This relay appears to be less than 2 weeks old. This [blog post](#) explains the lifecycle of a new relay, and why it will not be immediately fully used to capacity.

**Configuration**

**Nickname**  
pearl000

**OR Addresses**  
35.203.138.1:9001

**Contact**  
pearl@m15n.org <https://github.com/mmcloughlin/pearl>

**Dir Address**  
none

**Advertised Bandwidth**  
150 KiB/s

**IPv4 Exit Policy Summary**  
reject 1-65535

**IPv6 Exit Policy Summary**  
reject 1-65535

**Exit Policy**  
reject \*:\*

**Effective Family Members**

**Properties**

**Fingerprint**  
68325B4B1E177374B84D372F03046351BEE7FF6A

**Uptime**

**Flags**  
Running Valid

**Properties**  
none

**DNS Name**  
1.138.203.35.bc.googleusercontent.com

**Country**  
United States

**AS Number**  
AS15169

**AS Name**  
Google Inc.

**First Seen**  
2017-10-25 05:00:00 (6 days 11 hours 12 minutes and 25 seconds)

**Last Restarted**  
2017-10-25 05:00:00 (6 days 11 hours 12 minutes and 25 seconds)

**Consensus Weight**  
20

**Platform**  
Pearl cfa5d6 on Linux

# Future Work

There is *a long way* to go.

Near term:

- Nurse **production** deployment
- Develop more realistic local **testing**
- Focus on **performance**
- Revisit **parser**

Longer term expand protocol support:

- Exit node, client side, hidden services...

Alternatively:

- Contribute to new [go-tor project](#) or the [official tor project](#)

## Acknowledgements

- **Myles** for our daily check-in and keeping me on the straight and narrow
- **Oz** and **Myles** for setting up the program
- **Bradfield Sabbatical Program** for the community and emotional support
- **Uber** for letting me participate



# Thank you

Michael McLoughlin

Software Engineer, Uber

[mmcloughlin@gmail.com](mailto:mmcloughlin@gmail.com)

<http://mmcloughlin.com>

[@mbmcloughlin](#)